

Explainable Artificial Intelligence

VTR UGE 21

DAY 4 · MODULE 4

Reliability of Explanations

When can a probability, an explanation and a decision be trusted? Calibration, shift, manipulated explanations, fairness and a benchmark against random.

Prof. Dr. Utku Kose

Full Professor, Department of Computer Engineering, Süleyman Demirel University, Isparta, Türkiye
Founding Director, AI Application and Research Center (YAZEM), Süleyman Demirel University
Head of the Computer Science Division, Department of Computer Engineering, Süleyman Demirel University
Additional affiliations: University of North Dakota (USA), Universidad Panamericana (Mexico City, Mexico),
Vel Tech University (Chennai, India)
IEEE Senior Member, ACM Professional Member

ORCID 0000-0002-9652-6415 · utkukose.com · github.com/utkukose
utkukose@sdu.edu.tr · utku.kose@und.edu · ukose@up.edu.mx · utkukose@gmail.com

Overview

An explanation is only as reliable as the model behind it and the method that produced it. This day asks three questions. Can the probabilities of a model be taken at face value [1], and does the model still work when the population changes? Can an explanation be trusted, or can it be changed while the prediction stays the same [2, 3], and does it beat a random ranking [4]? Can the decisions be trusted, or do they treat groups differently through a hidden proxy, a feature that stands in for the group [5]? Each question comes with a test that can be run on any model.

Estimated study time: 6 to 8 hours.

Learning outcomes

By the end of the day, students are expected to read a reliability diagram and the expected calibration error, to detect a shifted population with a domain classifier, to explain why a model can be confidently wrong under shift, to describe how an explanation can be manipulated while the prediction stays the same, to audit a model for group differences and find a proxy with SHAP (Shapley additive explanations), and to test an explanation method against a random baseline. In Python, students are expected to build, filter and group pandas data frames and draw plots from them.

Python in this day

The Colab notebook opens with Python step 4: Data frames and plots, right after its setup section. It covers building a pandas data frame, summaries, boolean filtering, grouping, binning with `pd.cut` and plotting a reliability table, applied to a handful of credit applicants [6, 7]. After the step, the notebook continues with the full audits.

How these notes point to the Python code. Every method of this lecture is also written as Python code in the Colab notebook of the day, the file `NB04_reliability_of_explanations.ipynb`. A Colab notebook is a document of text and Python code that runs in the web browser, with nothing to install. It opens with the button Open in Colab on the day overview page of the course, <https://github.com/utkukose/xai-veltech-NB-lecture/tree/main/days/day-04>. The code of the notebook is divided into numbered sections. In these notes, each box titled *Python code in the Colab notebook, Section N* names the section whose code carries out what the text above the box explains, and says what to run and what to look at there. The animations that the notes mention are interactive figures on the lecture page of the day, <https://utkukose.github.io/xai-veltech-NB-lecture/days/day-04/lecture.html>. The interactive lab, with the practice parts and the self-assessment of the day, is at <https://utkukose.github.io/xai-veltech-NB-lecture/days/day-04/lab.html>.

Python step 4: Data frames and plots

Audits of models are mostly operations on tables: Filtering, grouping and binning. This step does each of them with pandas on a handful of credit applicants and draws a reliability table [6, 7].

A data frame and its summary

```
import pandas as pd
credit = pd.DataFrame({"income": [42000, 31000, 58000, 27000, 39000, 51000],
                      "late_payments": [0, 3, 1, 4, 2, 0],
                      "approved": [1, 0, 1, 0, 1, 1]})

print(credit)
print(credit.describe().round(1))
```

Output

	income	late_payments	approved
0	42000	0	1
1	31000	3	0
2	58000	1	1
3	27000	4	0
4	39000	2	1
5	51000	0	1

	income	late_payments	approved
count	6.0	6.0	6.0
mean	41333.3	1.7	0.7
std	11741.7	1.6	0.5
min	27000.0	0.0	0.0
25%	33000.0	0.2	0.2
50%	40500.0	1.5	1.0
75%	48750.0	2.8	1.0
max	58000.0	4.0	1.0

A data frame holds named columns of equal length. **describe** summarises every numerical column: The count, the mean, the standard deviation and the quartiles.

Filtering rows

```
late = credit[credit["late_payments"] >= 2]
print(late)
print(credit.loc[credit["income"] > 40000, ["late_payments", "approved"]])
```

Output

	income	late_payments	approved
1	31000	3	0
3	27000	4	0
4	39000	2	1

	late_payments	approved
0	0	1
2	1	1
5	0	1

A comparison with a column gives a column of booleans, and indexing the data frame with it keeps the rows where it is **True**. **loc** selects rows and columns at once, by condition and by name.

Grouping

```
credit["many_late"] = credit["late_payments"] >= 2
rates = credit.groupby("many_late")["approved"].mean()
print(rates)
print("difference of the approval shares:", round(float(rates.max() - rates.min()), 3))
```

Output

```
many_late
False    1.000000
True     0.333333
Name: approved, dtype: float64
difference of the approval shares: 0.667
```

A new column is created by assigning to a new name, as in a dictionary. **groupby** splits the rows by the values of a column, here whether an applicant has two or more late payments, and the mean of the approval column in each group is the share of approved applicants. All applicants with fewer late payments were approved, and one of the three with more.

A reliability table

```
probs = pd.Series([0.05, 0.12, 0.35, 0.41, 0.62, 0.68, 0.81, 0.93, 0.97, 0.22])
outcome = pd.Series([0, 0, 1, 0, 1, 1, 1, 1, 1, 0])
bins = pd.cut(probs, [0, 0.25, 0.5, 0.75, 1.0])
table = pd.DataFrame({"confidence": probs, "outcome": outcome}).groupby(bins,
observed=True).mean()
print(table)
```


Output

	confidence	outcome
(0.0, 0.25]	0.130000	0.0
(0.25, 0.5]	0.380000	0.5
(0.5, 0.75]	0.650000	1.0
(0.75, 1.0]	0.903333	1.0

pd.cut assigns every probability to an interval, and grouping by the intervals gives the mean confidence and the observed frequency of each bin, the two coordinates of a reliability diagram.

A plot from a data frame

```
import matplotlib.pyplot as plt
ax = table.plot(x="confidence", y="outcome", marker="o", legend=False, figsize=(4.5,
3.2))
ax.plot([0, 1], [0, 1], "k--", lw=1)
ax.set_xlabel("mean confidence"); ax.set_ylabel("observed frequency")
plt.show()
```



Output of the code above.

Data frames can plot themselves: `plot` returns the axes, on which further lines and labels can be drawn. The dashed diagonal marks perfect calibration.

Quick check. Why can the area under the receiver operating characteristic (ROC) curve not detect miscalibration?

Can the probability be trusted?

The diagnosis case so far

Days 1 to 3 explained the models of the diagnosis case. Readable models showed their reasons directly, the random forest of the tumour data was explained from the outside, and the pixel-level maps of a convolutional network pointed at the hidden marker of the lesion images in every test image.

The audit now turns from what the models do to how far their results can be trusted. The lesion network and the tumour data return in these tests, and a credit model joins them as a second case, because the treatment of groups of applicants raises the question of fairness.

The classifiers of this day return a probability. For a case with the input x , the model f returns $p = f(x)$, a number between 0 and 1: the predicted probability that the case belongs to the positive class, the class with the label $y = 1$. A model is calibrated when its probabilities mean what they say: Among cases scored 0.8, about 80 percent are positive. As a formula:

$$P(y = 1 \mid f(x) = p) = p \quad \text{for every probability } p$$

→ This formula with numbers, step by step: [Lab Part A](#), step 2.

In words: Among all cases to which the model gives the probability p , the share of positive cases is p . The day tests this on synthetic credit data, which the fairness section below describes in full. There are 4000 applicants with six features, and the positive class is that the applicant repays the loan, which 31 percent do. Every model is trained on 2800 applicants and judged on the other 1200, the test cases.

One exact probability value occurs too rarely to count a frequency, so the test cases are sorted into bins by their probability. The test uses ten bins of width 0.1: Bin 1 holds the cases with p below 0.1, bin 2 those from 0.1 to below 0.2, and so on. Let $p^{(i)}$ and $y^{(i)}$ be the probability and the label of test case i , and let n_b be the number of cases in bin b . Two averages describe the bin:

$$\bar{p}_b = \frac{1}{n_b} \sum_{i \in \text{bin } b} p^{(i)}, \quad \bar{y}_b = \frac{1}{n_b} \sum_{i \in \text{bin } b} y^{(i)}$$

→ These formulas with numbers, step by step: [Lab Part A](#), step 2.

The first average is the mean predicted probability of the bin, also called its confidence. The second is the observed frequency: Because the labels are 0 and 1, their mean is the share of positive cases. A reliability diagram plots the observed frequency against the mean predicted probability in bins, one point per bin. For a calibrated model the points lie on the diagonal.

The expected calibration error, abbreviated ECE, summarises the gaps [1]. It is the average distance of the points from the diagonal, in which every bin counts with its share n_b / n of the n test cases:

$$\text{ECE} = \sum_b \frac{n_b}{n} |\bar{y}_b - \bar{p}_b|$$

→ This formula with numbers, step by step: [Lab Part A](#), step 3.

The sum runs over the bins that contain at least one case. The absolute value keeps gaps above and below the diagonal from cancelling. An ECE of 0.05 says that the probabilities are off by five percentage points on average. Guo and colleagues define the error with the confidence in the predicted class and the accuracy of each bin. This course uses the form for two classes given here, with the probability and the frequency of the positive class.

In everyday life: a weather forecast

A forecaster who announces a 70 percent chance of rain is calibrated when it rains on about 7 of every 10 such days. Weather services check their forecasts in exactly this way, with the same kind of diagram. A forecaster who says 90 percent whenever rain is likely may still pick the rainy days well, and yet the numbers mislead everyone who plans by them.

Worked example: the calibration error of ten cases

Ten cases have the probabilities 0.05, 0.12, 0.35, 0.41, 0.62, 0.68, 0.81, 0.93, 0.97 and 0.22 and the labels 0, 0, 1, 0, 1, 1, 1, 1 and 0, the same cases as in Python step 4. Four bins of width 0.25 sort them. The first bin holds three cases with the probabilities 0.05, 0.12 and 0.22, none of them positive: Its mean probability is $(0.05 + 0.12 + 0.22) / 3 = 0.39 / 3 = 0.13$, its observed frequency is $0 / 3 = 0$, and its gap is $|0 - 0.13| = 0.13$. The second bin holds 0.35, a positive case, and 0.41, a negative one: Its mean is $(0.35 + 0.41) / 2 = 0.38$, its frequency $1 / 2 = 0.5$ and its gap $0.5 - 0.38 = 0.12$. The third bin holds 0.62 and 0.68, both positive: mean $1.30 / 2 = 0.65$, frequency 1, gap $1 - 0.65 = 0.35$. The fourth holds 0.81, 0.93 and 0.97, all positive: mean $2.71 / 3 = 0.903$, frequency 1, gap $1 - 0.903 = 0.097$.

The error weights every gap with the share of the cases in its bin, $3 / 10 = 0.3$, $2 / 10 = 0.2$, 0.2 and 0.3 : $0.3 \times 0.13 + 0.2 \times 0.12 + 0.2 \times 0.35 + 0.3 \times 0.097$, which is $0.039 + 0.024 + 0.070 + 0.029 = 0.162$.

Interactive lab, Part A.

The same ten cases with a choice of two, four, five or ten bins and of squared probabilities: the bins, the gaps, the calibration error and the Brier score, step by step, with the reliability diagram. Part A opens at <https://utkukose.github.io/xai-veltech-NB-lecture/days/day-04/lab.html#pa>.

The shape of the curve tells in which direction a model errs. A model is called overconfident when its probabilities are more extreme, closer to 0 and 1, than the observed frequencies justify, and underconfident when they stay too close to 0.5.

Shape of the curve	What the model does	Name
On the diagonal	Its probabilities match the observed frequencies	calibrated
Flatter than the diagonal: above it at low probabilities, below it at high ones	Its probabilities are more extreme than the frequencies	overconfident
Steeper than the diagonal: below it at low probabilities, above it at high ones	Its probabilities stay too close to 0.5	underconfident
Below the diagonal everywhere	Its probabilities are too high everywhere	overestimating
Above the diagonal everywhere	Its probabilities are too low everywhere	underestimating

Two common scores do not show calibration. The predicted class $\hat{y}$ is positive when the probability is above 0.5, and accuracy is the share of test cases whose predicted class equals the label:

$$\hat{y} = 1 \text{ if } p > 0.5, \quad \hat{y} = 0 \text{ otherwise}$$

$$\text{accuracy} = \frac{\text{number of test cases with } \hat{y}^{(i)} = y^{(i)}}{n}$$

Accuracy only asks on which side of 0.5 a probability lies: A probability of 0.51 counts like one of 0.99. The second score is the area under the receiver operating characteristic curve, abbreviated ROC curve. The area is abbreviated AUC, for area under the curve. It equals the probability that a randomly chosen positive case receives a higher probability from the model than a randomly chosen negative case:

$$\text{AUC} = P(p^{(i)} > p^{(k)} \mid y^{(i)} = 1, y^{(k)} = 0)$$

$$= \frac{\text{number of correctly ordered pairs}}{n_1 \times n_0}$$

Here n_1 and n_0 are the numbers of positive and negative test cases. A pair consists of one positive case i and one negative case k . It is correctly ordered when the positive case has the higher probability, and a pair with equal probabilities counts as half. An AUC of 0.5 is the level of guessing, and 1 is a perfect order. The ROC curve itself plots, for every possible threshold, the true positive rate against the false positive rate, two rates that the fairness section below defines. The AUC depends only on the order of the probabilities. Any change that keeps the order, such as squaring every probability, leaves it unchanged, although it changes the calibration.

A third score does react to calibration. The Brier score is the mean squared distance between the probability and the label:

$$\text{Brier} = \frac{1}{n} \sum_{i=1}^n (p^{(i)} - y^{(i)})^2$$

→ This formula with numbers, step by step: [Lab Part A](#), step 4.

It is 0 for a model that gives every positive case the probability 1 and every negative case the probability 0, and lower is better. A useful yardstick is the model that gives every applicant the same probability, the

share of repayers 0.31. Each repayer contributes $(1 - 0.31)^2 = 0.476$ and each other applicant $0.31^2 = 0.096$, so its Brier score is $0.31 \times 0.476 + 0.69 \times 0.096 = 0.148 + 0.066 = 0.214$, the same as 0.31×0.69 .

Worked example: the same order, another calibration

Six of the ten cases of the first worked example are positive and four are negative, which makes $6 \times 4 = 24$ pairs. In 23 of them the positive case has the higher probability. The exception is the positive case with 0.35 and the negative case with 0.41. The AUC is $23 / 24 = 0.958$. The ten squared distances are 0.0025, 0.0144, 0.4225, 0.1681, 0.1444, 0.1024, 0.0361, 0.0049, 0.0009 and 0.0484, for example $(0.35 - 1)^2 = 0.4225$ for the positive case with 0.35. They add up to 0.9446, so the Brier score is $0.9446 / 10 = 0.094$.

Squaring every probability keeps the order, so the AUC stays at 0.958. The calibration becomes worse. The squared probabilities fall into bins of 5, 2, 1 and 2 cases with the gaps 0.129, 0.577, 0.344 and 0.097, so the calibration error rises from 0.162 to $0.5 \times 0.129 + 0.2 \times 0.577 + 0.1 \times 0.344 + 0.2 \times 0.097 = 0.234$, and the Brier score from 0.094 to 0.161. The accuracy changes as well, from 0.9 to 0.7, because the probabilities 0.62 and 0.68 fall below 0.5.

Interactive lab, Part B.

Four reliability diagrams, each to be named: calibrated, overconfident, underconfident or too high everywhere.

Part B opens at <https://utkukose.github.io/xai-veltech-NB-lecture/days/day-04/lab.html#pb>.

A miscalibrated model can be repaired after training. Recalibration learns a calibration map c that turns each probability p into a corrected probability $c(p)$. The simplest map is binning, which the animation below uses: Every probability is replaced by the observed frequency of its bin, measured on cases that were set aside for this purpose:

$$c(p) = \bar{y}_b \quad \text{for every } p \text{ in bin } b$$

The day uses isotonic regression, a finer version of the same idea. Isotonic means non-decreasing: The map never goes down when the probability goes up, so the order of the cases is never reversed. Among all such maps, the method chooses the one whose corrected probabilities lie closest to the labels:

$$\min_{c \text{ non-decreasing}} \sum_i (y^{(i)} - c(p^{(i)}))^2$$

The result is a staircase whose steps can have any width. The map has to be fitted on cases that the model did not see during training. The training data are therefore split into five parts. Five times, the model is trained on four parts and the map is fitted on the fifth, and the final probability is the average of the five corrected models.

An ECE computed from 1200 test cases is itself uncertain. The bootstrap measures this uncertainty. It draws 1200 cases at random from the test cases, where the same case may be drawn several times, and computes the ECE of this resample. Repeating this 200 times gives 200 values. The 5th and the 95th percentile, the values below which 5 and 95 percent of them fall, form the bootstrap range. A change of the ECE counts as real only when the ranges before and after do not overlap.

Three credit models are compared. Gradient boosting builds its prediction from many small decision trees, each of which corrects the errors of the trees before it. The default version uses 100 trees of depth 3. The deep version uses 400 trees of depth 6 and fits the 2800 training applicants without a single error. Naive Bayes is a simple model that treats the features as independent within each class. Some models are

well calibrated out of the box; others, such as deep ensembles of trees, are overconfident and need recalibration.




Figure. Reliability diagrams of three credit models on the 1200 test applicants, before recalibration (raw) and after isotonic recalibration. The dashed diagonal marks perfect calibration.

In the figure, the red raw curve of the default model follows the diagonal, and only its last three bins, which hold 24, 10 and 2 applicants, scatter. Its ECE is 0.009. The deep model is overconfident. It gives 671 of the 1200 test applicants a probability below 0.1, yet 22 percent of them repay, and of the 95 applicants scored above 0.9 only 53 percent repay. Its curve is much flatter than the diagonal, its ECE is 0.203, and its Brier score of 0.249 is worse than the 0.214 of the constant probability. Isotonic recalibration brings the blue curve back to the diagonal.

Model	ECE, raw (bootstrap range)	ECE, recalibrated (bootstrap range)	Brier score, raw and recalibrated	AUC, raw and recalibrated
gradient boosting	0.009 (0.017 to 0.040)	0.019 (0.020 to 0.046)	0.186 and 0.187	0.721 and 0.721
deep gradient boosting	0.203 (0.186 to 0.227)	0.026 (0.019 to 0.050)	0.249 and 0.198	0.660 and 0.668
naive Bayes	0.027 (0.028 to 0.057)	0.019 (0.017 to 0.047)	0.194 and 0.191	0.703 and 0.703

Recalibration repairs the deep model: Its bootstrap ranges before and after, 0.186 to 0.227 and 0.019 to 0.050, lie far apart. For the other two models the ranges overlap, so the change is noise, and the ECE of the default model even rises from 0.009 to 0.019. The AUC hardly moves in any row, because recalibration hardly changes the order of the applicants. The table holds one more lesson: The ECE of the default model, 0.009, lies below its own bootstrap range. A resample adds chance differences in every bin, and absolute gaps cannot cancel, so chance tends to raise the measured error. For the same reason a perfectly calibrated model does not reach an ECE of 0 on a finite sample.

The animation produces miscalibration on purpose. It starts from simulated cases whose true probability p is known and reports a distorted probability $\tilde{p}$. The distortion works on the log-odds $\ln(p / (1 - p))$, also called the logit of p . It multiplies the log-odds by an overconfidence factor a and turns the result back into a probability with the logistic function σ :

$$\tilde{p} = \sigma\left(a \cdot \ln \frac{p}{1 - p}\right), \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$

With $a = 1$ nothing changes. With $a = 2$ a true probability of 0.7 has the log-odds $\ln(0.7 / 0.3) = \ln 2.333 = 0.847$, which become $2 \times 0.847 = 1.694$, and it is reported as $1 / (1 + e^{-1.694}) = 1 / 1.184 = 0.84$. A probability of 0.3 has the log-odds -0.847 and is reported as $1 / (1 + e^{1.694}) = 1 / 6.441 = 0.16$. The model is overconfident. With a below 1 the probabilities move towards 0.5, and the model is underconfident. A

probability of 0.5 stays at 0.5 for every factor, so no case changes its side and the accuracy stays the same. In the animation the accuracy is 0.736 for every factor, while the ECE grows from 0.027 at $a = 1$ to 0.108 at $a = 2$ and 0.183 at $a = 4$. Recalibration by binning measures the frequencies on one half of the 4000 cases and corrects the other half, which brings the ECE at $a = 2$ down to 0.024.

Animation on the lecture page: A reliability diagram of 2000 simulated cases. Raise the overconfidence factor and watch the curve bend away from the diagonal while the accuracy stays the same. Then recalibrate by binning.

Python code in the Colab notebook, Section 1.

Open Section 1 of the Colab notebook and run it. The function `make_credit` simulates the applicants, and the function `ece` is the calibration error above, computed bin by bin. The cell measures the AUC, the Brier score and the ECE of three models, draws their reliability diagrams and recalibrates them with `CalibratedClassifierCV(method="isotonic", cv=5)`. The bootstrap in `boot` gives the interval of each ECE.

Check your understanding (1). Two models have the same accuracy, but one is overconfident. What differs?

- A. Nothing
- B. How well their probabilities match the observed frequencies
- C. Their ranking of the cases
- D. Their training data

Check your understanding (2). A model is tested on 200 cases. Half of them fall into a bin with a mean predicted probability of 0.2 and an observed frequency of 0.3. The other half fall into a bin with a mean predicted probability of 0.9 and an observed frequency of 0.7. What is the expected calibration error?

- A. 0.05
- B. 0.15
- C. 0.20
- D. 0.30

Has the population changed?

A model is validated on data like its training data. When the population changes, for example after an economic downturn, both the model and its explanations can become wrong. Such a change between training and use is called distribution shift. A simple test trains a classifier to tell the old data from the new, a domain classifier. If it cannot do better than chance, the two samples look alike.

The test needs no labels of the original problem. Every case of the old data receives the new label 0 and every case of the new data the label 1, and a classifier is trained to predict this label from the features. Here the classifier is a random forest: 200 decision trees, each trained on a random resample of the data, whose votes are averaged. Every case is scored by a forest that has not seen it. For this purpose the data are split into five parts, and each part is scored by a forest trained on the other four, which is called five-fold cross-validation. The AUC of these scores measures the shift: A value near 0.5 means that old and new cases cannot be told apart, and a value near 1 means that they hardly overlap.

The new data come from a simulated downturn: 1500 new applicants drawn from the generator of the credit data, with every income lowered by 18 percent, one late payment added to every applicant and every debt ratio raised by 15 percent. The domain classifier separates these applicants from the 2800 training applicants with an AUC of 0.725, a clear shift. The credit model suffers as well: Its AUC falls from 0.721 on the original test data to 0.694 on the new population.

In everyday life: a rule learned in another season

A commuter who knows from summer that the drive to work takes 25 minutes is late on the first snowy day. Nothing in the rule itself says that the conditions have changed. Only a look at today, compared with the days on which the rule was learned, shows it, and a domain classifier makes this comparison for a model.

The domain classifier also tells which features changed. A random forest reports an importance for every feature, computed from the Gini impurity of the nodes of its trees. A node in which a share ρ of the cases belongs to class 1 has the impurity $2\rho(1 - \rho)$, the Gini impurity of Day 1: 0 when all cases belong to one class and 0.5 when the classes are mixed half and half. A split is useful when the two child nodes are purer than their parent. Its value is the decrease of impurity, weighted by the share of all n cases that reach the node:

$$\Delta = \frac{n_{\text{node}}}{n} \left[\text{Gini}_{\text{node}} - \frac{n_{\text{left}}}{n_{\text{node}}} \text{Gini}_{\text{left}} - \frac{n_{\text{right}}}{n_{\text{node}}} \text{Gini}_{\text{right}} \right]$$

The impurity importance of a feature adds up Δ over all splits that use the feature. The sums are averaged over the trees and scaled so that the importances of all features add up to 1. For example, a node with 100 cases, half of them of class 1, has the impurity $2 \times 0.5 \times 0.5 = 0.5$. A split into two nodes of 50 cases with 40 and 10 cases of class 1 gives the children the shares $40 / 50 = 0.8$ and $10 / 50 = 0.2$ and both the impurity $2 \times 0.8 \times 0.2 = 0.32$. When this node is the root, with $n_{\text{node}} = n$, the decrease is $0.5 - 0.5 \times 0.32 - 0.5 \times 0.32 = 0.5 - 0.32 = 0.18$.

Feature	Mean in the training data	Mean in the new data	Importance in the domain classifier
income	33899	27960	0.199
debt ratio	0.300	0.351	0.186
postcode score	53.5	53.3	0.164
late payments	1.52	2.53	0.158
years of credit history	6.65	6.44	0.149
years of employment	4.97	4.97	0.145

The importances must be read with care. Income and the debt ratio rank first, and both did change. The number of late payments changed most clearly, yet it ranks fourth, behind the postcode score, which did not change at all. The cause lies in the measure: Impurity importance favours features with many different values, which offer many places to split. Income takes thousands of different values and the number of late payments about ten. Without any shift, a forest of this kind gives the late payments an importance of only 0.065 and each of the other five features between 0.17 and 0.20. The shift therefore shows in the rise of the late payments to 0.158, not in the first places of the ranking. A comparison of the two samples feature by feature, as in the table, has to confirm what the importances suggest.

Python code in the Colab notebook, Section 2.

Open Section 2 of the Colab notebook and run it. The cell simulates the downturn, trains the domain classifier with `cross_val_predict` and prints its AUC, `roc_auc_score(yd, p_dom)`. Its attribute `feature_importances_` shows which features have changed.

The second experiment of this section shows how badly a shift can end. Day 1 called a cue that predicts the label in the training data without being its reason a shortcut [8]. The lesion images of Days 1 and 3 contain one. An image has 32 by 32 pixels with grey values between 0 and 1. Class 1 has a lesion, a faint round blob near the centre of the image, and class 0 has none. The marker is a bright square of 4 by 4 pixels in the top left corner. In the training data it is added to 97 percent of the images with a lesion and to 3 percent of the images without one. The small convolutional network of Day 3, with about 10000 weights, is trained again on 2400 such images.

A slice is a part of the test data chosen by a rule. Three slices are built from the same 600 new images, which differ only in the marker. In the first slice the marker follows the rule of the training data. In the second it is added to half of the images of either class, so it carries no information. In the third the rule is reversed: The marker is added to 3 percent of the images with a lesion and to 97 percent of those without.

For every image x the network computes two scores, $s_0(x)$ and $s_1(x)$, one per class. Such scores are called logits. The softmax function turns them into two probabilities that add up to 1:

$$P(y = k | x) = \frac{e^{s_k(x)}}{e^{s_0(x)} + e^{s_1(x)}}, \quad k = 0, 1$$

The predicted class is the class with the larger probability, and this probability is called the confidence of the prediction:

$$\hat{y} = \arg \max_k P(y = k | x), \quad \text{confidence} = \max_k P(y = k | x)$$

With two classes the confidence lies between 0.5 and 1. Two numbers describe each slice: the accuracy and the mean confidence, the average of the confidence over the 600 images.

Worked example: confidence from two scores

For one image the network returns the scores $s_0 = 0.2$ and $s_1 = 3.2$. With $e^{0.2} = 1.22$ and $e^{3.2} = 24.53$, the probability of class 1 is $24.53 / (1.22 + 24.53) = 24.53 / 25.75 = 0.953$ and that of class 0 is $1.22 / 25.75 = 0.047$. The predicted class is 1, with a confidence of 0.953.

The confidence depends only on the difference of the two scores, here 3. It is computed from the image alone, without the label, so it cannot know whether the prediction is right.

Slice	Images in which the marker follows the rule of the training data	Accuracy	Mean confidence
same distribution	95.8 percent	0.958	0.973
marker uninformative	48.2 percent	0.482	0.964
marker reversed	3.0 percent	0.030	0.955

A model that learned a shortcut is the extreme case: The lesion network of Day 3 is about 96 percent accurate on familiar images, near 50 percent when the marker carries no information and about 3 percent when the marker convention is reversed, with a confidence above 95 percent throughout. The second

column of the table explains the pattern. In every image of every slice the network predicts a lesion exactly when the marker is present. Its accuracy therefore equals the share of images in which the marker follows the rule of the training data. The confidence does not react, because the network is certain about the marker, and the marker is still there. The gap between mean confidence and accuracy, the calibration error in its simplest form, grows from $0.973 - 0.958 = 0.015$ on the first slice to $0.955 - 0.030 = 0.925$ on the third.

Python code in the Colab notebook, Section 3.

Open Section 3 of the Colab notebook and run it. `make_lesion_images` builds the three slices, each with its own `marker_rate`, and `softmax_rows` turns the outputs of the network into the probabilities `P`. The cell prints for each slice the accuracy, from `P.argmax(1)`, and the mean confidence, from `P.max(1)`.

Check your understanding (3). On the reversed slice, the network is wrong 97 percent of the time with a confidence of 95 percent. What does this show?

- A. The network is broken
- B. Confidence measures the model's certainty about its shortcut, not its correctness
- C. The labels are wrong
- D. Calibration fixes it

Can the explanation be manipulated?

Gradient-based explanations depend on the local shape of the model, and that shape can change sharply between points with almost the same prediction [2]. Dombrowski and colleagues showed that an explanation can be changed almost at will while the output stays fixed [3].

The explanation of this section is gradient times input [9]. Let $s_t(x)$ be the score that a network gives to its predicted class t for the input $x = (x_1, \dots, x_d)$ with d features. The gradient $\partial s_t / \partial x_j$ is the slope of the score in the direction of feature j : It says how much the score changes when the feature grows by a small amount. The attribution of feature j multiplies this slope by the value of the feature:

$$\phi_j(x) = x_j \cdot \frac{\partial s_t(x)}{\partial x_j}, \quad j = 1, \dots, d$$

→ This formula with numbers, step by step: [Lab Part C](#), step 4.

A positive ϕ_j means that the feature, at its present value, supports the predicted class, and a negative one that it speaks against it. The d attributions together form the explanation $\phi(x)$. Two explanations are compared with the cosine similarity. For two lists of attributions ϕ and ψ , the Greek letter psi, it is

$$\cos(\phi, \psi) = \frac{\sum_{j=1}^d \phi_j \psi_j}{\sqrt{\sum_{j=1}^d \phi_j^2} \sqrt{\sum_{j=1}^d \psi_j^2}}$$

→ This formula with numbers, step by step: [Lab Part C](#), step 5.

The numerator multiplies the attributions feature by feature and adds the products. The two roots are the lengths of the lists, so the result does not depend on their overall size. It lies between -1 and 1 : 1 when the two explanations point the same way, 0 when they have nothing in common and -1 when they are

opposite. For example, the explanations (3, 4) and (4, 3) have the products $3 \times 4 + 4 \times 3 = 24$ and the lengths $\sqrt{(3^2 + 4^2)} = 5$, so their cosine similarity is $24 / (5 \times 5) = 0.96$.

The animation shows the mechanism in two dimensions. Its model has two features, called x and y in the animation and x_1 and x_2 here, and the output

$$f(x_1, x_2) = \sigma(3(\sin(2x_1) + x_2))$$

→ This formula with numbers, step by step: [Lab Part C](#), step 2.

On the curve $x_2 = 0.4 - \sin(2x_1)$ the inner bracket is always 0.4, so the prediction is $\sigma(3 \times 0.4) = \sigma(1.2) = 1 / (1 + e^{-1.2}) = 1 / 1.301 = 0.769$, more exactly 0.7685, at every point of the curve. The gradient is not the same at every point. By the chain rule of Day 3, with f for the prediction,

$$\frac{\partial f}{\partial x_1} = 6f(1-f)\cos(2x_1), \quad \frac{\partial f}{\partial x_2} = 3f(1-f)$$

→ These formulas with numbers, step by step: [Lab Part C](#), step 3.

On the curve the second slope is $3 \times 0.7685 \times 0.2315 = 0.534$ everywhere, while the first follows $\cos(2x_1)$ and even changes its sign. The animation applies gradient times input to these two slopes.

Worked example: one prediction, two explanations

At the start of the curve, $x_1 = -1.2$ and $x_2 = 0.4 - \sin(-2.4) = 0.4 + 0.675 = 1.075$. With $6f(1-f) = 1.0674$ and $\cos(-2.4) = -0.7374$ the gradient is $(1.0674 \times (-0.7374), 0.534) = (-0.787, 0.534)$, and gradient times input gives the explanation $((-1.2) \times (-0.787), 1.075 \times 0.534) = (0.944, 0.574)$: Both features support the prediction.

At $x_1 = 1.0$ the curve passes through $x_2 = 0.4 - \sin 2 = 0.4 - 0.909 = -0.509$, and the prediction is still 0.769. With $\cos 2 = -0.4161$ the gradient is $(1.0674 \times (-0.4161), 0.534) = (-0.444, 0.534)$ and the explanation $(1.0 \times (-0.444), -0.509 \times 0.534) = (-0.444, -0.272)$: Now both features speak against the prediction.

The products of the two explanations add up to $0.944 \times (-0.444) + 0.574 \times (-0.272) = -0.419 - 0.156 = -0.575$. The lengths are $\sqrt{(0.944^2 + 0.574^2)} = 1.105$ and $\sqrt{(0.444^2 + 0.272^2)} = 0.521$, and the cosine similarity is $-0.575 / (1.105 \times 0.521) = -0.575 / 0.576 = -1.00$.

Interactive lab, Part C.

Two points of the same curve chosen with sliders: their gradients, their explanations and the cosine similarity of the two, step by step, with the explanations drawn as arrows. Part C opens at <https://utkukose.github.io/xai-veltech-NB-lecture/days/day-04/lab.html#pc>.

In everyday life: the same mark, different working

Two students can write the same correct answer under an exam question, one with a sound derivation and one with a lucky guess dressed up as a derivation. The mark is the same and the working differs. The prediction of a model is the mark, and an explanation is the working, which can be changed while the mark stays.

Animation on the lecture page: Points along a curve of equal prediction. Move along the curve and watch the explanation, gradient times input, change while the prediction stays at 0.7685.

The attack of the day targets a small neural network on the tumour data of the week [10]: 569 tumours, each described by 30 measurements of cell nuclei and labelled malignant, class 0, or benign, class 1. Every measurement is standardised, that is, shifted and scaled so that on the training data its mean is 0 and its standard deviation is 1, so one unit of a standardised feature is one standard deviation. The network has 30 inputs, one hidden layer of 48 units and two outputs, and it classifies 97.9 percent of the 143 test tumours correctly.

The attack searches for a perturbation δ : a list of 30 small changes, one per measurement. The perturbed input $x + \delta$ should receive an explanation that is as dissimilar to the original one as possible, while the probability p_t of the predicted class stays where it was. Both wishes are written into one function, which the search makes as small as it can:

$$L(\delta) = \cos(\phi(x + \delta), \phi(x)) + \lambda (p_t(x + \delta) - p_t(x))^2$$

The first term is the cosine similarity between the new and the original explanation. The second term punishes any change of the probability, and the constant λ sets how strongly: It is 60 here. The size of the perturbation is limited by a budget ε . No measurement may change by more than $\varepsilon = 0.35$ standard deviations, $|\delta_j| \leq \varepsilon$ for every feature j .

The search is gradient descent. It starts from a tiny random perturbation with values between -0.01 and 0.01 , because at $\delta = 0$ the function L is flat and gives no direction. In each of 400 steps it computes the gradient G of L with respect to δ , moves against it and cuts every δ_j back into the range from $-\varepsilon$ to ε , which the function clip does. The step uses Adam, short for adaptive moment estimation, the method that also trains the networks of Day 3. Adam keeps a running average m of the gradient and a running average v of its square, and divides one by the root of the other:

$$\begin{aligned} m &\leftarrow 0.9m + 0.1G, & v &\leftarrow 0.999v + 0.001G^2 \\ \delta &\leftarrow \text{clip}\left(\delta - \eta \frac{\hat{m}}{\sqrt{\hat{v}} + 10^{-8}}, -\varepsilon, \varepsilon\right) \end{aligned}$$

The constant η , the Greek letter eta, is the learning rate, 0.02 here. All operations act on each of the 30 coordinates separately, and m and v start at zero. The hats mark a correction for this start: In step k , $\hat{m} = m / (1 - 0.9^k)$ and $\hat{v} = v / (1 - 0.999^k)$. Because of the division, a coordinate moves by about 0.02 per step, whatever the size of its gradient.

The attack targets the explanation of the fourth test patient, a benign tumour that the network recognises with a probability of 0.9945. The search ends with 29 of the 30 measurements at the edge of the budget. Within 0.35 standard deviations per measurement, the explanation of one patient changes substantially while the predicted probability moves by less than one percentage point, from 0.9945 to 0.9871. The cosine similarity between the explanations before and after is 0.69.




Figure. The ten largest attributions of one tumour case before the attack, and the attributions of the same features after it.

The figure shows the ten features with the largest attributions before the attack. Mean compactness remains a strong supporter of the prediction, but its attribution falls from 1.78 to 1.07. Worst area falls from 0.50 to almost zero, from rank 3 to rank 30, the last. Area error grows from 0.29 to 1.08 and becomes the top feature, and radius error grows from -0.30 to -0.89 . The eight most important features of the original explanation move by 7.8 ranks on average, and only four of them are still among the top eight. A reviewer shown the second explanation would draw different conclusions about the same patient with the same diagnosis. An explanation shown to a decision maker therefore needs a robustness check of its own.

Python code in the Colab notebook, Section 4.

Open Section 4 of the Colab notebook and run it. The function `explanation` returns the gradient times input of the network, and the function `attack` searches for a small change `d` of the patient, limited by `eps`, that changes the explanation while the prediction stays. The cell compares the prediction and the explanation before and after the attack.

Check your understanding (4). Why can an explanation change strongly while the prediction stays the same?

- A. The model is linear
- B. Explanations depend on local derivatives, which can differ between points with the same prediction
- C. The labels changed
- D. The data shifted

Can the decision be trusted? Fairness and proxies

The credit data of the day are synthetic, so everything about them is known. Each of the 4000 applicants belongs to group A, 65 percent, or to group B, 35 percent. The group stands for a protected attribute: a property such as sex, origin or religion on which a decision must not depend. Six features describe an applicant, and the label says whether the loan is repaid. The label is drawn at random with the probability $\sigma(z)$, in which the log-odds z add up the features with fixed weights:

$$\begin{aligned}
z = & 0.9 (\ln \text{income} - 10.4) - 2.4 \times \text{debt ratio} \\
& + 0.045 \times \text{history years} - 0.30 \times \text{late payments} \\
& + 0.035 \times \text{employment years} + 1.2 \times \frac{\text{postcode score} - 60}{40}
\end{aligned}$$

The last term ties the outcome to the home district of the applicant, and the postcode score differs strongly between the groups, as the table shows. In the whole data, 36.8 percent of group A and 20.3 percent of group B repay.

Feature	Meaning	Mean in group A	Mean in group B
income	income of the applicant	36375	29756
debt_ratio	debt as a share of income, between 0 and 1	0.286	0.328
history_years	length of the credit history in years	6.52	6.71
late_payments	number of late payments	1.34	1.83
employment_years	years in employment	5.01	4.75
postcode_score	score of the home district, between 0 and 100	59.9	42.3

The default gradient boosting model of the first section is trained on the six features. An applicant counts as approved, $\hat{y} = 1$, when the predicted probability of repaying is above 0.5. The credit model never sees the group of an applicant, yet it approves 22 percent of group A and 3 percent of group B. A postcode score that is correlated with the group carries the difference: Removing the protected attribute does not remove a disparity when a proxy remains. A proxy is a feature that stands in for the protected attribute because the two are correlated.

Two common measures summarise such differences [5]. Both compare rates between the groups. The approval rate of a group is the share of its applicants who are approved, written $P(\hat{y} = 1 \mid A)$ for group A. Two further rates look at a part of the applicants only:

$$\text{TPR} = P(\hat{y} = 1 \mid y = 1), \quad \text{FPR} = P(\hat{y} = 1 \mid y = 0)$$

→ These formulas with numbers, step by step: [Lab Part D](#), steps 1 and 2.

The true positive rate, abbreviated TPR, is the share of approved applicants among those who repay. The false positive rate, abbreviated FPR, is the share of approved applicants among those who do not. Demographic parity compares approval rates; equal opportunity compares approval rates among those who would repay, that is, true positive rates:

$$\Delta_{\text{DP}} = |P(\hat{y} = 1 \mid A) - P(\hat{y} = 1 \mid B)|$$

$$\Delta_{\text{EO}} = |P(\hat{y} = 1 \mid y = 1, A) - P(\hat{y} = 1 \mid y = 1, B)|$$

→ These formulas with numbers, step by step: [Lab Part D](#), steps 3 and 4.

The first line is the demographic parity difference and the second the equal opportunity difference. Each is 0 when the two groups are treated alike by its measure.

Worked example: two fairness measures

A model decides about 200 applicants of group A and 150 of group B. In group A, 100 applicants would repay. The model approves 70 of them and 20 of the 100 who would not. The approval rate is $(70 + 20) / 200 = 90 / 200 = 0.45$, the true positive rate is $70 / 100 = 0.70$, and the false positive rate is $20 / 100 = 0.20$.

In group B, 60 applicants would repay. The model approves 27 of them and 9 of the other 90. The approval rate is $(27 + 9) / 150 = 36 / 150 = 0.24$, the true positive rate is $27 / 60 = 0.45$, and the false positive rate is $9 / 90 = 0.10$.

The demographic parity difference is $0.45 - 0.24 = 0.21$, and the equal opportunity difference is $0.70 - 0.45 = 0.25$.

Interactive lab, Part D.

The same two groups with sliders for the approved applicants of each kind: the three rates of both groups and the two differences, step by step. Part D opens at <https://utkukose.github.io/xai-veltech-NB-lecture/days/day-04/lab.html#pd>.

On the test data, 172 of the 793 applicants of group A are approved and 12 of the 407 of group B. The approval rates are $172 / 793 = 0.2169$ and $12 / 407 = 0.0295$, and the demographic parity difference is $0.2169 - 0.0295 = 0.187$. Among those who repay, 103 of 298 are approved in group A and 7 of 75 in group B: The true positive rates are $103 / 298 = 0.3456$ and $7 / 75 = 0.0933$, and the equal opportunity difference is $0.3456 - 0.0933 = 0.252$.

Interactive lab, Part E.

Two groups of 100 applicants with other counts: both fairness measures computed by hand and checked. Part E opens at <https://utkukose.github.io/xai-veltech-NB-lecture/days/day-04/lab.html#pe>.


How closely the postcode score follows the group is measured by the correlation coefficient r . For two columns of numbers $u_1, \dots, u_n$ and $v_1, \dots, v_n$ with the means $\bar{u}$ and $\bar{v}$, it is

$$r = \frac{\sum_{i=1}^n (u_i - \bar{u})(v_i - \bar{v})}{\sqrt{\sum_{i=1}^n (u_i - \bar{u})^2} \sqrt{\sum_{i=1}^n (v_i - \bar{v})^2}}$$

The coefficient lies between -1 and 1 . It is positive when the two columns rise together, negative when one falls as the other rises, and 0 when they are unrelated. Between the postcode score and membership of group B, coded as 1 for group B and 0 for group A, it is -0.57 : Applicants of group B live in districts with lower scores.

In everyday life: what a postcode reveals

A lender may not use the origin of an applicant, and yet the postcode can carry much of the same information, because people of similar background often live in the same districts. Deleting the forbidden column does not delete what the remaining columns reveal about it. An audit therefore looks for the feature that carries the gap, as the next paragraphs do.



Interactive scene on the lecture page: What a postcode reveals. A small town of sixty applicants in three districts with different postcode scores. Group B, in orange, lives mostly in the district with the lowest score. The credit score adds two terms of the generating formula of this section, the income and the postcode score, never the group, and approves the best quarter of the town. Removing the postcode score, and splitting the gap between the groups into one part per term, show how much of the difference the proxy carries.

The audit then asks which feature carries the gap. It uses SHAP, short for Shapley additive explanations, the method of Day 2 [11]. The SHAP value ϕ_j of feature j is its Shapley value: the change of the prediction when the feature joins a set S of other features, averaged over all subsets of the set F of the d features:

$$\phi_j = \sum_{S \subseteq F \setminus \{j\}} \frac{|S|! (d - |S| - 1)!}{d!} [v(S \cup \{j\}) - v(S)]$$

Here $v(S)$ is the prediction of the model when only the features in S are known, $|S|$ is the number of features in S , and the exclamation mark is the factorial, $3! = 3 \times 2 \times 1 = 6$. The fraction is the weight of a subset. The values of one case add up to its prediction, counted from a base value ϕ_0 , the average prediction. For a gradient boosting model the prediction is explained on the scale of the log-odds. For applicant i with the predicted probability $p^{(i)}$:

$$\ln \frac{p^{(i)}}{1 - p^{(i)}} = \phi_0 + \sum_{j=1}^d \phi_j^{(i)}$$

Averaging SHAP values by group shows which feature carries the gap. For every feature, the audit subtracts the mean SHAP value of the n_B test applicants of group B from that of the n_A of group A:

$$\Delta_j = \frac{1}{n_A} \sum_{i \in A} \phi_j^{(i)} - \frac{1}{n_B} \sum_{i \in B} \phi_j^{(i)}$$

Because the values of every applicant add up to the log-odds, the gaps Δ_j of the six features add up to the difference between the mean log-odds of the two groups.

Feature	Mean SHAP value in group A	Mean SHAP value in group B	Gap
postcode score	0.158	-0.232	0.390
income	0.086	-0.133	0.219
late payments	0.035	-0.080	0.114
debt ratio	0.027	-0.076	0.103
years of credit history	-0.010	0.004	-0.014
years of employment	-0.004	-0.002	-0.002

The gaps add up to $0.390 + 0.219 + 0.114 + 0.103 - 0.014 - 0.002 = 0.810$: The mean log-odds are -0.63 in group A and -1.44 in group B, and $-0.63 - (-1.44) = 0.81$. The postcode score carries 0.39 of this difference, $0.39 / 0.81 = 0.48$, almost half, although the group itself is not among the features. The rest comes from income, late payments and the debt ratio, which differ between the groups in the generating process. The audit ends with a test of its own finding: The model is trained again without the postcode score.

Measure	With the postcode score	Without the postcode score
AUC on the test data	0.721	0.679
Demographic parity difference	0.187	0.120
Equal opportunity difference	0.252	0.112

Both differences shrink, at a price in the quality of the ranking. They do not vanish, because the other features differ between the groups as well. Whether the remaining differences may be used is a question of policy, which an explanation can inform but not decide. Explanations can even be fooled into hiding a proxy [12].

Python code in the Colab notebook, Section 5.

Open Section 5 of the Colab notebook and run it. The function `group_metrics` computes the approval rate and the true and false positive rates of each group. The cell then explains the model with `shap.TreeExplainer`, compares the SHAP values by group to find the proxy `postcode_score`, and trains a second model, `gb2`, without it.

Check your understanding (5). Why does removing the protected attribute from the inputs not remove a disparity?

- A. The model is biased by design
- B. Other features correlated with the attribute can carry the same information
- C. SHAP is wrong
- D. The groups are equal

Does the explanation beat random?

A simple test checks whether an explanation knows anything about the model. The deletion test removes the features that the explanation ranks highest for the predicted class first and watches the probability of

that class fall. The insertion test adds them first to a neutral case. A faithful explanation makes the probability fall and rise quickly, and every method must at least beat a random ranking [4].

The test of the day runs on a random forest of 200 trees for the tumour data, which classifies about 96 percent of the test tumours correctly. A feature cannot be taken out of a table, so removing it means replacing its value by a neutral one: the median of the feature in the training data, the value below which half of the tumours lie. The neutral case x' is the row of the 30 medians. For a test case x , let t be the class that the forest predicts. The attributions are computed for class 1, the benign class. A feature therefore supports the predicted class with ϕ_j when $t = 1$ and with $-\phi_j$ when $t = 0$, and the features are ordered from the strongest support downwards.

The deletion curve starts with the probability q_0 that the forest gives to class t for the unchanged case. Then the features are replaced by their medians one at a time in this order, and q_k is the probability of class t after k replacements. The insertion curve runs the other way: It starts at the neutral case, and step k gives the feature of rank k its real value. The test changes the first $K = 15$ of the 30 features. Each curve is summarised by the area under it, computed with trapezoids between neighbouring points and divided by K . This is the mean height of the curve:

$$\text{Del} = \frac{1}{K} \sum_{k=1}^K \frac{q_{k-1} + q_k}{2}, \quad \text{faithfulness} = \text{Ins} - \text{Del}$$

The area Del belongs to the deletion curve, and Ins is the same average over the points of the insertion curve. A good explanation has a small deletion area and a large insertion area, and their difference serves as one score of faithfulness. The [Colab notebook](#) and the [interactive lab](#) call the two areas deletion AUC and insertion AUC. AUC again means area under a curve, but here the curve is a deletion or insertion curve, not the ROC curve of the first section.

Worked example: two areas and a faithfulness score

A deletion curve over $K = 4$ features has the probabilities 0.90, 0.60, 0.40, 0.30 and 0.25. The four trapezoids have the mean heights $(0.90 + 0.60) / 2 = 0.75$, $(0.60 + 0.40) / 2 = 0.50$, $(0.40 + 0.30) / 2 = 0.35$ and $(0.30 + 0.25) / 2 = 0.275$, so the area is $(0.75 + 0.50 + 0.35 + 0.275) / 4 = 1.875 / 4 = 0.469$.

The insertion curve of the same explanation has the probabilities 0.25, 0.55, 0.75, 0.85 and 0.90, with the mean heights $(0.25 + 0.55) / 2 = 0.40$, 0.65, 0.80 and 0.875 and the area $(0.40 + 0.65 + 0.80 + 0.875) / 4 = 2.725 / 4 = 0.681$. The faithfulness is $0.681 - 0.469 = 0.212$.

Four rankings are compared on 20 test tumours. The first uses SHAP values, computed with the exact method for trees of Day 2. The second uses LIME, short for local interpretable model-agnostic explanations [13], in a short version of its own. LIME draws 600 artificial cases z around the case x by adding random noise to every feature, asks the forest for their probabilities $f(z)$ and fits a linear surrogate model g to them:

$$g(z) = \beta_0 + \sum_{j=1}^d \beta_j u_j, \quad u_j = \frac{z_j - x_j}{\text{sd}_j}$$

Here u_j is the distance of the artificial case from the real one in feature j , counted in standard deviations sd_j of the training data, and the noise has the size of one standard deviation. The coefficients β_j are chosen so that g follows the forest closely, most of all near the case:

$$\min_{\beta} \sum_z \pi_x(z) (f(z) - g(z))^2 + \alpha \sum_{j=1}^d \beta_j^2$$

The first sum runs over the 600 artificial cases. The weight $\pi_x(z)$ is large for artificial cases close to x :

$$\pi_x(z) = \exp\left(-\frac{u_1^2 + \dots + u_d^2}{\kappa^2}\right), \quad \kappa = 0.75 \sqrt{d}$$

The constant κ is the kernel width: It sets how fast the weight falls with the distance. The value $0.75\sqrt{d}$ is the default of the LIME library, 4.1 for the 30 features. The last term of the fit, with $\alpha = 1$, keeps the coefficients small. The coefficients β_j are the attributions of LIME. The third ranking is global and is labelled impurity (global) in the figure: It orders the features by the impurity importance of the forest, the measure of the second section, which is the same for every case and has no sign. The fourth ranking orders the features by random numbers.

Ranking	Deletion area, lower is better	Insertion area, higher is better	Faithfulness
SHAP	0.733	0.921	0.188
global ranking	0.766	0.893	0.127
LIME	0.818	0.856	0.038
random ranking	0.881	0.781	-0.101

On the tumour data SHAP is the most faithful method of the comparison [11], and every method beats the random ranking. The random ranking does not score 0 but -0.101 , because its deletion curve stays high and its insertion curve stays low. This value, not zero, is the mark to beat. LIME with the default kernel width beats it clearly, yet it stays behind the global ranking, which ignores the case altogether. As Day 2 showed, the result of LIME depends on the kernel width. A benchmark like this one, run before a method is chosen, replaces the assumption that a popular method is a reliable one.

In everyday life: beating a benchmark that costs nothing

A fund manager is judged against an index fund that anyone can buy, not against zero. An explanation method is judged in the same way against a ranking that costs nothing, the random one. Here the random ranking scores -0.101 , and that, not zero, is the mark to beat.




Figure. Mean deletion and insertion curves of four methods on 20 cases of the tumour data.

The figure shows the mean curves of the 20 cases. All deletion curves start at 0.955, the mean probability of the predicted class for the unchanged cases. After one replacement the SHAP curve has fallen to 0.894, while the random curve has not moved. After 15 replacements the curves stand at 0.649 for SHAP, 0.697 for the global ranking, 0.744 for LIME and 0.807 for the random ranking. The insertion curves start at 0.699 and reach 0.982, 0.948, 0.921 and 0.860 in the same order.

The deletion curves cannot fall to zero, for a reason that lies in the neutral case. The forest calls the row of medians benign with a probability of 0.997. Of the 20 cases, 14 are predicted benign, and replacing their measurements by benign-looking medians hardly lowers their probability: With the SHAP order it falls from 0.96 to 0.92. The six malignant cases decide the test. With the SHAP order their probability of being malignant falls from 0.94 to 0.34 after five replacements and to 0.02 after fifteen. With the random order it only falls to 0.43. The choice of the neutral case is part of the test.

Part F of the interactive lab runs the same test in the browser. Its model has eight features and the output $f(x) = \sigma(w_1 x_1 + \dots + w_8 x_8 + x_1 x_2)$, with the weights 2.0, -1.5, 1.2, 0.8, -0.5, 0.3, 0.1 and 0. The neutral value of every feature is 0, the local explanation is gradient times input, and all eight features are changed, so $K = 8$. When all features are changed, every deletion curve ends at the neutral case and every insertion curve at the real case, whatever the order. A good order then shows in a deletion curve that dips below its end value, because the features that speak against the prediction are removed last.

Interactive lab, Part F.

The deletion and insertion tests in the browser on a model of eight features, for gradient times input, a random order and other rankings, with a scoreboard of the areas. Part F opens at <https://utkukose.github.io/xai-veltech-NB-lecture/days/day-04/lab.html#pf>.

Python code in the Colab notebook, Section 6.

Open Section 6 of the Colab notebook and run it. The function `curve` changes the features of a patient one after another in the order of an attribution and records the prediction, and `np.trapezoid` gives the area under the curve. The cell reports the deletion area, the insertion area and their difference for SHAP, LIME, a global ranking and a random ranking.

Check your understanding (6). What must every explanation method beat in a deletion test?

- A. SHAP
- B. A random ranking
- C. LIME
- D. A perfect ranking

Check your understanding (7). A deletion curve over three features has the probabilities 0.8, 0.5, 0.3 and 0.2. The insertion curve of the same explanation has the probabilities 0.2, 0.6, 0.8 and 0.8. What is the faithfulness, the insertion area minus the deletion area?

- A. 0.10
- B. 0.20
- C. 0.43
- D. 0.63

The diagnosis case, stage 4: how far the results can be trusted

The methods of this day now come together in the fourth stage of the project. The audit no longer asks what the models do but how far their results can be trusted: their probabilities, their behaviour on a


changed population, their explanations and their decisions about groups of people. The first figure shows the stage of the project, the second the five tests of the day.

The diagnosis case, stage 4 of 5



Figure. The stage of the project reached today: the results of the audit tested for trust.

The diagnosis case after Day 4: five tests of how far the results can be trusted



Result of stage 4: Probabilities can be repaired, a shift and a proxy can be found, and an explanation needs a check of its own against manipulation and against chance.

Figure. The audit after Day 4: five tests of trust on the lesion network, the tumour data and the credit model.

Question	Answer of stage 4	Section
Do the probabilities mean what they say?	For the default credit model, with an ECE of 0.009. The deep model is overconfident, with 0.203, and isotonic recalibration brings it to 0.026	Can the probability be trusted?
Has the population changed?	A domain classifier separates a simulated downturn with an AUC of 0.725, and the AUC of the credit model falls from 0.721 to 0.694	Has the population changed?
Does the confidence warn of a shortcut?	No: The lesion network stays above 0.95 while its accuracy falls from 0.958 to 0.030	Has the population changed?
Can an explanation be manipulated?	Yes: Changes of at most 0.35 standard deviations turn the explanation of a tumour to a cosine similarity of 0.69, while the probability moves from 0.9945 to 0.9871	Can the explanation be manipulated?
Are the decisions fair?	The postcode score carries 0.39 of the gap of 0.81 in the log-odds, and without it the equal opportunity difference falls from 0.252 to 0.112	Can the decision be trusted? Fairness and proxies
Does the explanation beat random?	SHAP reaches a faithfulness of 0.188, against -0.101 for a random ranking	Does the explanation beat random?

The audit now holds its evidence: what the models rely on, where they look and how far their results can be trusted. Day 5 turns this evidence into artefacts for the people who use the models: a record of every decision, a dashboard, a model card and the checks that regulation asks for.

Reliability in your field

The application of the day uses the five fields of Days 1 and 2. For the model of each field it draws the reliability diagram and runs the deletion test with SHAP against a random order. Each field is a yes or no problem, and a default gradient boosting model is trained on 60 percent of the cases and tested on the other 40 percent. The table lists the fields with their results.

Field	Cases, features and positive class	Test accuracy	ECE	Deletion area: SHAP order and random order
medicine	569 tumours with 30 measurements; benign tumour	0.952	0.039	0.625 and 0.796
chemistry	178 wines with 13 chemical measurements; wine of the first of three cultivars	0.944	0.056	0.744 and 0.774
manufacturing	1500 simulated weld seams with 5 process values; sound weld	0.778	0.048	0.477 and 0.689
energy	1500 simulated transformers with 5 gas concentrations in the oil; fault in the transformer	0.737	0.044	0.467 and 0.676
civil engineering	1500 simulated concrete mixes with 5 properties; strength above 30 megapascals	0.897	0.045	0.498 and 0.817

The deletion test of this section differs from the benchmark in three settings. It uses the first 15 test cases, it replaces all d features and not only 15, and it divides the area by d . Because all features are replaced, the SHAP curve and the random curve end at the same value, the probability of the neutral case, and the area only measures how early a curve falls. In every field the SHAP order has the smaller area. The reliability diagram of the tumour model looks jagged: Of its 228 test tumours, 220 are scored below 0.1 or above 0.9, and the bins in between hold one or two tumours each. The ECE of 0.039 is small all the same, because these bins carry little weight.

Python code in the Colab notebook, Section 7.

Open Section 7 of the Colab notebook. Its first line is a switch, `FIELD = "medicine"`. Replace the word by `chemistry`, `manufacturing`, `energy` or `civil` and run the section. The function `load_field` loads or simulates the data of that field, and the cell trains a model on them, measures its calibration error with the function `ece` of Section 1 and runs the deletion test of SHAP against a random ranking with `audit_curve`.

Going further (optional)

The attack needs the derivative of the explanation itself, which is a derivative of a derivative. For a network with one smooth hidden layer it can be written as a matrix, and this section checks it against finite differences. The network computes the score of its predicted class in two stages. Each of the 48 hidden units h forms a weighted sum a_h of the 30 inputs, and the score adds up the hidden units after a smooth function, the softplus $\ln(1 + e^a)$:

$$s_t(x) = b_t + \sum_{h=1}^{48} v_h \ln(1 + e^{a_h}), \quad a_h = b_h + \sum_{j=1}^{30} w_{jh} x_j$$

Here w_{jh} is the weight from input j to hidden unit h , v_h is the weight from hidden unit h to the score, and b_h and b_t are constants called biases. The softplus is a smooth version of $\max(0, a)$, the function that the networks of Day 3 use. Its slope is the logistic function $\sigma(a)$, and the slope of $\sigma(a)$ is $\sigma(a)(1 - \sigma(a))$. The chain rule gives the gradient and the second derivative of the score:

$$\frac{\partial s_t}{\partial x_j} = \sum_{h=1}^{48} v_h \sigma(a_h) w_{jh}$$

$$\frac{\partial^2 s_t}{\partial x_j \partial x_l} = \sum_{h=1}^{48} v_h \sigma(a_h) (1 - \sigma(a_h)) w_{jh} w_{lh}$$

The second line is a matrix of 30 by 30 numbers, one for every pair of features j and l . It says how the slope for feature j changes when feature l moves. The matrix is useful only because the softplus is smooth. With $\max(0, a)$ the slopes would jump at some inputs and stay constant in between, so the second derivative would say nothing about how they change. The attribution $\phi_j = x_j \partial s_t / \partial x_j$ then changes with feature l as follows:

$$\frac{\partial \phi_j}{\partial x_l} = x_j \frac{\partial^2 s_t}{\partial x_j \partial x_l} + \begin{cases} \partial s_t / \partial x_j & \text{if } l = j \\ 0 & \text{otherwise} \end{cases}$$

The gradient of the attack function L follows from these pieces by the chain rule, once through the cosine similarity and once through the probability. Such a derivation is easy to get wrong, so it is checked with numbers. A finite difference estimates a slope from two values of the function, one a small step τ , the Greek letter tau, above the point and one the same step below it:

$$\frac{\partial L}{\partial \delta_j} \approx \frac{L(\dots, \delta_j + \tau, \dots) - L(\dots, \delta_j - \tau, \dots)}{2\tau}, \quad \tau = 10^{-6}$$

The dots stand for the other coordinates of the perturbation, which keep their values. For example, the function x^3 has the slope $3 \times 2^2 = 12$ at $x = 2$, and the finite difference with $\tau = 0.001$ gives $(2.001^3 - 1.999^3) / 0.002 = 12.000001$. For the network of the day the formula and the finite differences agree for all 30 features: Their largest difference, divided by the largest finite difference, is 5.3×10^{-9} .

A second experiment varies the budget of the attack: The larger the allowed change of the input, the further the explanation can be pushed, while the prediction barely moves. With a budget of half a standard deviation the cosine similarity falls to 0.30, and in no run does the probability change by more than 0.0074.

Budget in standard deviations	Cosine similarity after the attack	Change of the predicted probability
0.05	0.994	0.0001
0.10	0.973	0.0022

Budget in standard deviations	Cosine similarity after the attack	Change of the predicted probability
0.20	0.890	0.0010
0.35	0.691	0.0074
0.50	0.303	0.0025

Python code in the Colab notebook, Section 8.

Open Section 8 of the Colab notebook, which is optional, and run it. The function `attack_loss` holds the loss of the attack and its derivative, written out by hand, and the cell compares this derivative with finite differences for all 30 features. It then repeats the attack for five budgets `eps` and reports the damage.

Answers to the knowledge checks

Quick check. Because it depends only on the ranking of the scores: Changing the scores in a monotone way, which changes calibration, leaves the ranking and the area unchanged.

- 1. B.** Calibration concerns the meaning of the probabilities, not the order of the cases.
- 2. B.** Each bin holds half of the cases, and the gaps are $|0.3 - 0.2| = 0.1$ and $|0.7 - 0.9| = 0.2$: $0.5 \times 0.1 + 0.5 \times 0.2 = 0.05 + 0.10 = 0.15$. The curve is flatter than the diagonal, the mark of an overconfident model.
- 3. B.** A model can be confidently wrong when the data change in the way its shortcut depends on.
- 4. B.** Equal outputs do not imply equal gradients.
- 5. B.** A proxy reproduces the attribute's information.
- 6. B.** A method that cannot beat random adds no information about the model.
- 7. B.** The mean heights of the deletion curve are $(0.8 + 0.5) / 2 = 0.65$, 0.40 and 0.25 , so its area is $(0.65 + 0.40 + 0.25) / 3 = 1.30 / 3 = 0.433$. The insertion curve gives $(0.40 + 0.70 + 0.80) / 3 = 1.90 / 3 = 0.633$, and the difference is $0.633 - 0.433 = 0.20$.

References

- [1] Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017). On calibration of modern neural networks. In Proceedings of the 34th International Conference on Machine Learning (ICML 2017), PMLR 70, 1321-1330. <https://proceedings.mlr.press/v70/guo17a.html>
- [2] Ghorbani, A., Abid, A., & Zou, J. (2019). Interpretation of neural networks is fragile. In Proceedings of the AAAI Conference on Artificial Intelligence, 33(01), 3681-3688.
- [3] Dombrowski, A.-K., Alber, M., Anders, C. J., Ackermann, M., Müller, K.-R., & Kessel, P. (2019). Explanations can be manipulated and geometry is to blame. In Advances in Neural Information Processing Systems 32 (NeurIPS 2019).
- [4] Hedström, A., Weber, L., Krakowczyk, D., Bareeva, D., Motzkus, F., Samek, W., Lapuschkin, S., & Höhne, M. M.-C. (2023). Quantus: An explainable AI toolkit for responsible evaluation of neural network explanations and beyond. Journal of Machine Learning Research, 24(34), 1-11. <https://jmlr.org/papers/v24/22-0142.html>
- [5] Barocas, S., Hardt, M., & Narayanan, A. (2023). Fairness and Machine Learning: Limitations and Opportunities. MIT Press. <https://fairmlbook.org/>

- [6] McKinney, W. (2010). Data structures for statistical computing in Python. In Proceedings of the 9th Python in Science Conference (pp. 56-61). <https://doi.org/10.25080/Majora-92bf1922-00a>
- [7] Hunter, J. D. (2007). Matplotlib: A 2D graphics environment. Computing in Science & Engineering, 9(3), 90-95.
- [8] Geirhos, R., Jacobsen, J.-H., Michaelis, C., Zemel, R., Brendel, W., Bethge, M., & Wichmann, F. A. (2020). Shortcut learning in deep neural networks. Nature Machine Intelligence, 2(11), 665-673. <https://doi.org/10.1038/s42256-020-00257-z>
- [9] Ancona, M., Ceolini, E., Öztireli, C., & Gross, M. (2018). Towards better understanding of gradient-based attribution methods for deep neural networks. In 6th International Conference on Learning Representations (ICLR 2018). arXiv:1711.06104.
- [10] Wolberg, W., Mangasarian, O., Street, N., & Street, W. (1993). Breast Cancer Wisconsin (Diagnostic) [Dataset]. UCI Machine Learning Repository. <https://doi.org/10.24432/C5DW2B>
- [11] Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. In Advances in Neural Information Processing Systems 30 (NeurIPS 2017) (pp. 4765-4774). <https://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions>
- [12] Slack, D., Hilgard, S., Jia, E., Singh, S., & Lakkaraju, H. (2020). Fooling LIME and SHAP: Adversarial attacks on post hoc explanation methods. In Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society (AIES '20) (pp. 180-186). ACM. <https://doi.org/10.1145/3375627.3375830>
- [13] Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). "Why should I trust you?" Explaining the predictions of any classifier. In Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '16) (pp. 1135-1144). ACM.

Prof. Dr. Utku Kose

Full Professor, Department of Computer Engineering, Süleyman Demirel University, Isparta, Türkiye

Founding Director, AI Application and Research Center (YAZEM), Süleyman Demirel University

Head of the Computer Science Division, Department of Computer Engineering, Süleyman Demirel University

Additional affiliations: University of North Dakota (USA), Universidad Panamericana (Mexico City, Mexico), Vel Tech University (Chennai, India)

IEEE Senior Member, ACM Professional Member

ORCID 0000-0002-9652-6415 · utkukose.com · github.com/utkukose

utkukose@sdu.edu.tr · utku.kose@und.edu · ukose@up.edu.mx · utkukose@gmail.com